

Open Source | Open Possibilities



## Code size reduction using Similar Function Merging

Tobias Edler von Koch  
(University of Edinburgh,  
Qualcomm Innovation Center)

Pranav Bhandarkar  
(Qualcomm Innovation Center)




# Outline

1. Why optimize for code size?
2. The Problem of Duplicate Code
3. Existing LLVM MergeFunctions Pass
4. Similar Function Merging
5. Results

# Why optimize for code size?

- Traditionally three goals of compiler optimization:
  - **Performance**
  - **Power**
  - **Code size**
- External factors determine relative importance; there are complex interactions.
- Code size is key in many embedded scenarios



1 MB  $\approx 1/16^{\text{th}}$  of Cortex-A8 die size

# Code Size Reduction Approaches

Three main types:

- Hardware-based, e.g. ARM Thumb ISA.
- Software-based:
  - By **re-tuning** standard optimizations, e.g. inlining thresholds, loop unroll factor, etc.
  - By **actively reducing** code size of existing user code, e.g. elimination of redundancy.

We're looking at the last category.

# The Problem of Duplicate Code

- Software contains duplicate code due to:
  - ① Laziness, a.k.a. copy & paste
  - ② Manual templating
  - ③ C++ templates
  - ④ Compiler optimizations
- It may be possible for the user to fix ① & ② but ③ and ④ are much harder to control
- All types of duplication occur across the board in SPEC benchmarks, embedded systems code, ...

# Example from 400.perlbench

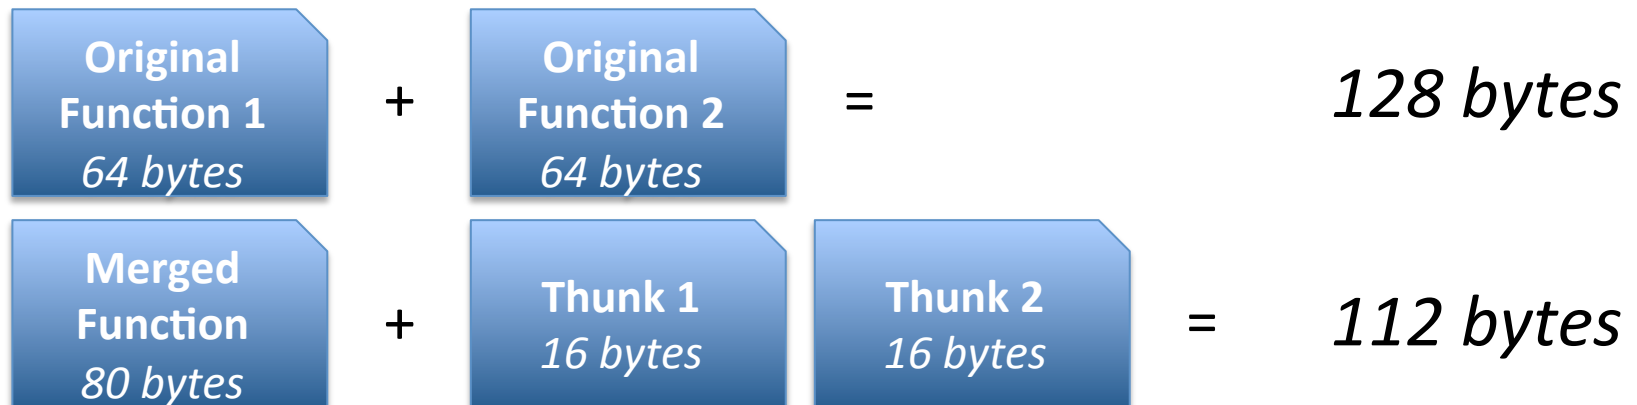
```
1 | OP *Perl_scalarkids(pTHX_ OP *o) {  
2 |     OP *kid;  
3 |     if (o && o->op_flags & OPf_KIDS) {  
4 |         for (kid = cLISTOPo->op_first; kid; kid = kid->op_sibling)  
5 |             scalar(kid);  
6 |     }  
7 |     return o;  
8 | }
```

```
1 | OP *Perl_listkids(pTHX_ OP *o) {  
2 |     OP *kid;  
3 |     if (o && o->op_flags & OPf_KIDS) {  
4 |         for (kid = cLISTOPo->op_first; kid; kid = kid->op_sibling)  
5 |             list(kid);  
6 |     }  
7 |     return o;  
8 | }
```

**Only a 1-instruction difference between  
the two functions in LLVM IR!**

# Example from 400.perlbench

- Merge the two functions:
  - Combine code from both in a new ‘merged function’
  - Insert if-statement where there are differences
  - Replace original functions with calls to merged function
- In our case, on x86:



Total savings: **12.5%**

# Existing MergeFunctions Pass

- Pass originally written by Nick Lewycky
- Disabled by default
- Merges ‘identical’ functions
- Introduces two key concepts we rely on:
  - Notion of **structural similarity** of functions to make analysis tractable
  - **Pointer-pointer-integer equivalence**:  
*pointers and integers of the same size* are treated as equivalent.... except where the difference matters.
- What if functions aren’t quite identical? We should still be able to merge them!



# Structural Similarity

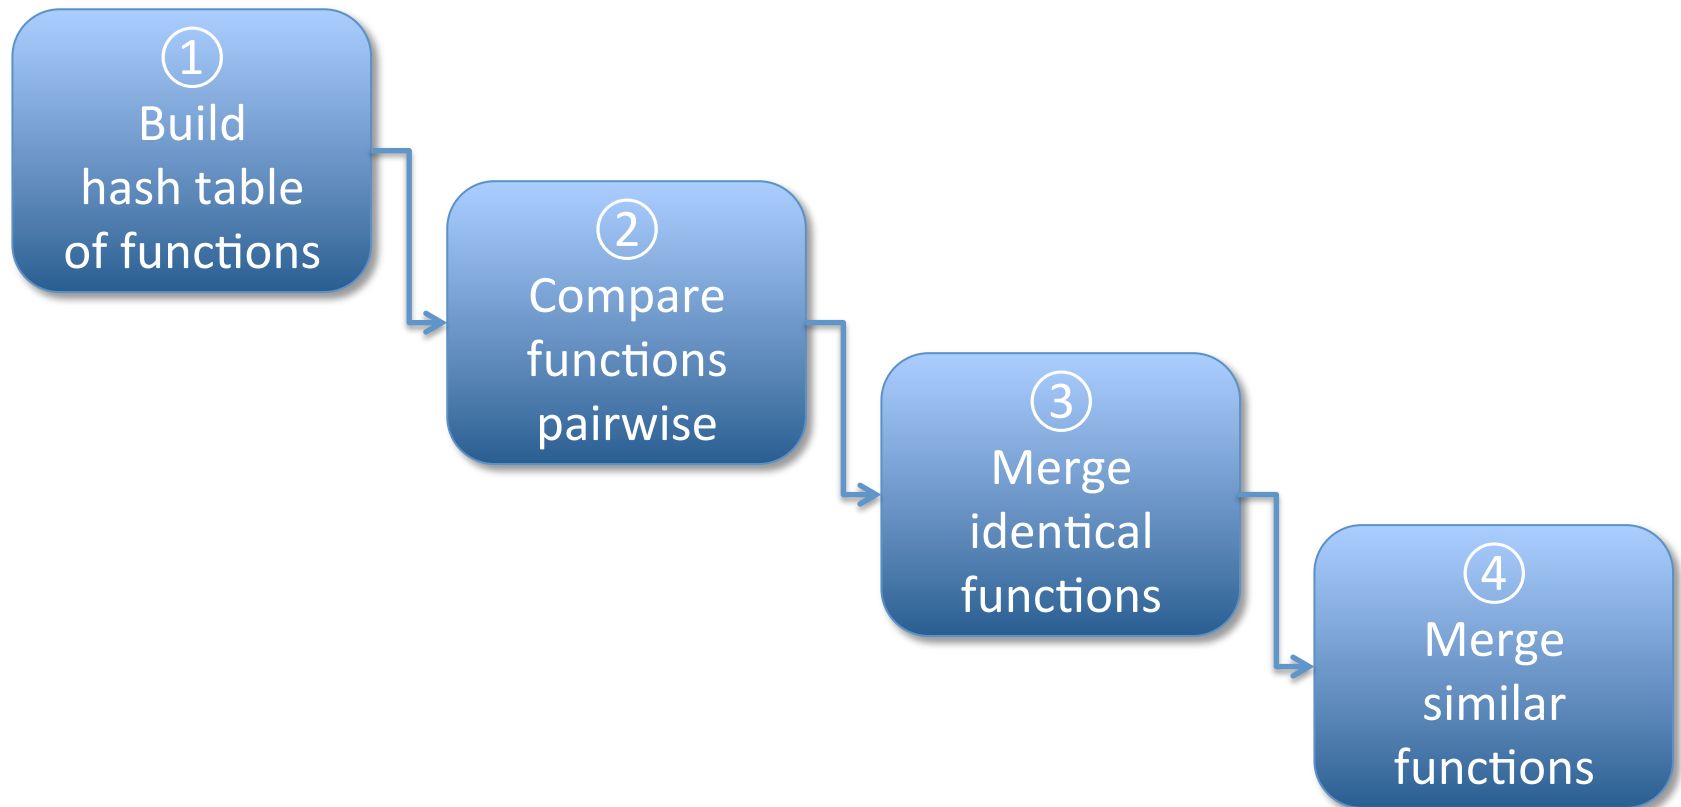
- Comparing all functions would be  $O(n^2)$   
... and we could theoretically merge everything!
- Introduce a number of practical constraints:

Functions must have

- Equivalent control flow graph and signature
- Same number of instructions in corresponding basic blocks  
**but:** allow differences in what these instructions are
- A minimum amount of similarity

# Similar Function Merging

The algorithm involves four main steps:




# Similar Function Merging

- **Step 1:** Insert functions into a hash table
  - Based on signature, number of basic blocks, ...
  - This avoids comparing functions that have no chance of being merged anyway
- **Step 2:** Compare all functions in each bucket

Still  $O(n^2)$  worst case, but better in practice

  - Follow control flow and compare block-by-block, instruction-by-instruction
  - Mark differing instructions
  - Give up if control flow or basic block length differs

# Example from 400.perlbench




# Similar Function Merging

- **Step 3:** Merge **identical** functions
  - Update call sites after merging
  - Other functions may become more similar as a result
  - Re-compare functions that have changed


Iterate this process until a **fixed point** is reached

# Similar Function Merging


- **Step 4:** Merge similar functions
  - Order pairs of functions by similarity
  - Pick *most similar* pair (A, B)
  - Find all (A, B') for which there is not a (B', C) with greater similarity
  - Merge A with B and the B's
  - Remove all pairs involving A, B, and the B's
  - Repeat this until there are no more functions to merge



Set of similar functions



Merged functions



# Similar Function Merging

- Run as a late optimization
- Tricky bits I haven't mentioned:
  - Must maintain SSA form throughout
  - Have to compare, update, and insert PHINodes:  
you can't put a conditional around two differing PHINodes
  - Thresholds are ISA-specific, need tuning for each arch
- How well does it work?

# Results

- We run the pass on
  - SPEC CPU2006 (Integer & FP benchmarks)
    - x86
    - Qualcomm Krait <sup>™</sup> (ARMv7-A)
    - Hexagon DSP <sup>™</sup>
  - A significant embedded application at QuIC
- Using LLVM/Clang 3.3



# Results

- To be added

# Conclusions

- Function merging is a promising technique for code size reduction
- Can reduce total code size for SPEC benchmarks by over 4% on x86
- We need a stronger focus on code size optimizations
  - as LLVM adoption in the embedded world increases this is becoming more critical

A panoramic banner image of the Edinburgh skyline, showing historic buildings and the city's architecture under a clear sky.

# EuroLLVM

Edinburgh, Scotland



April 7-8, 2014

# Thank you

and see you in Edinburgh!