

A Journey of OpenCL 2.0 Development in Clang

Anastasia Stulova
anastasia.stulova@arm.com
Media Processing Group
ARM

Agenda

- OpenCL intro
- OpenCL in Clang
- Overview of OpenCL 2.0
- OpenCL 2.0 implementation
- Summary and discussions

OpenCL programming model and terminology

C + OpenCL API

Host:


- Creates application
- Cross compiles for Device
- Sends work to Device
- Copy data to/from Device global memory

Offload




OpenCL Language

Device



- WI – work-item is a single sequential unit of work with its private memory (PM)
- WG – work-group is a collection of WIs that can run in parallel and access local memory shared among all WIs in the same WG

OpenCL intro




- C99 based
- Parallel units of work – kernels
- Explicitly assign object to memory using address space qualifier with each type
- Special types: images, events, pipes, ...
- Access qualifiers - read/write only applies to some types
- No standard C includes or libs, but defines its own libs

OpenCL for compiler writer

Ideal

How we imagined it to be




Reality

What we missed

- How to handle invalid targets?
- Conflicts between C and OpenCL unresolved in Spec (especially in undefined behaviour)!
- How to generate IR generically with absence of enough info on various backends?
- Missing explicit IR constructs are substituted with metadata and intrinsics!

First implementation in Clang (OpenCL 1.1/1.2)



OpenCL 2.0 feature overview

- Hierarchical/Dynamic parallelism - device side enqueue (work creation bypassing host) using ObjC blocks
- Reduce difficulty of writing code with address spaces (abstract away from memory model as much as possible, late binding)
- Simplify communications among kernels (avoid going outside of device via host)
 - Program scope variables persist across kernel invocations
 - Pipe communication using streaming pattern
- C11 atomics with memory visibility scope
- New image types and access qualifier

Generic address space

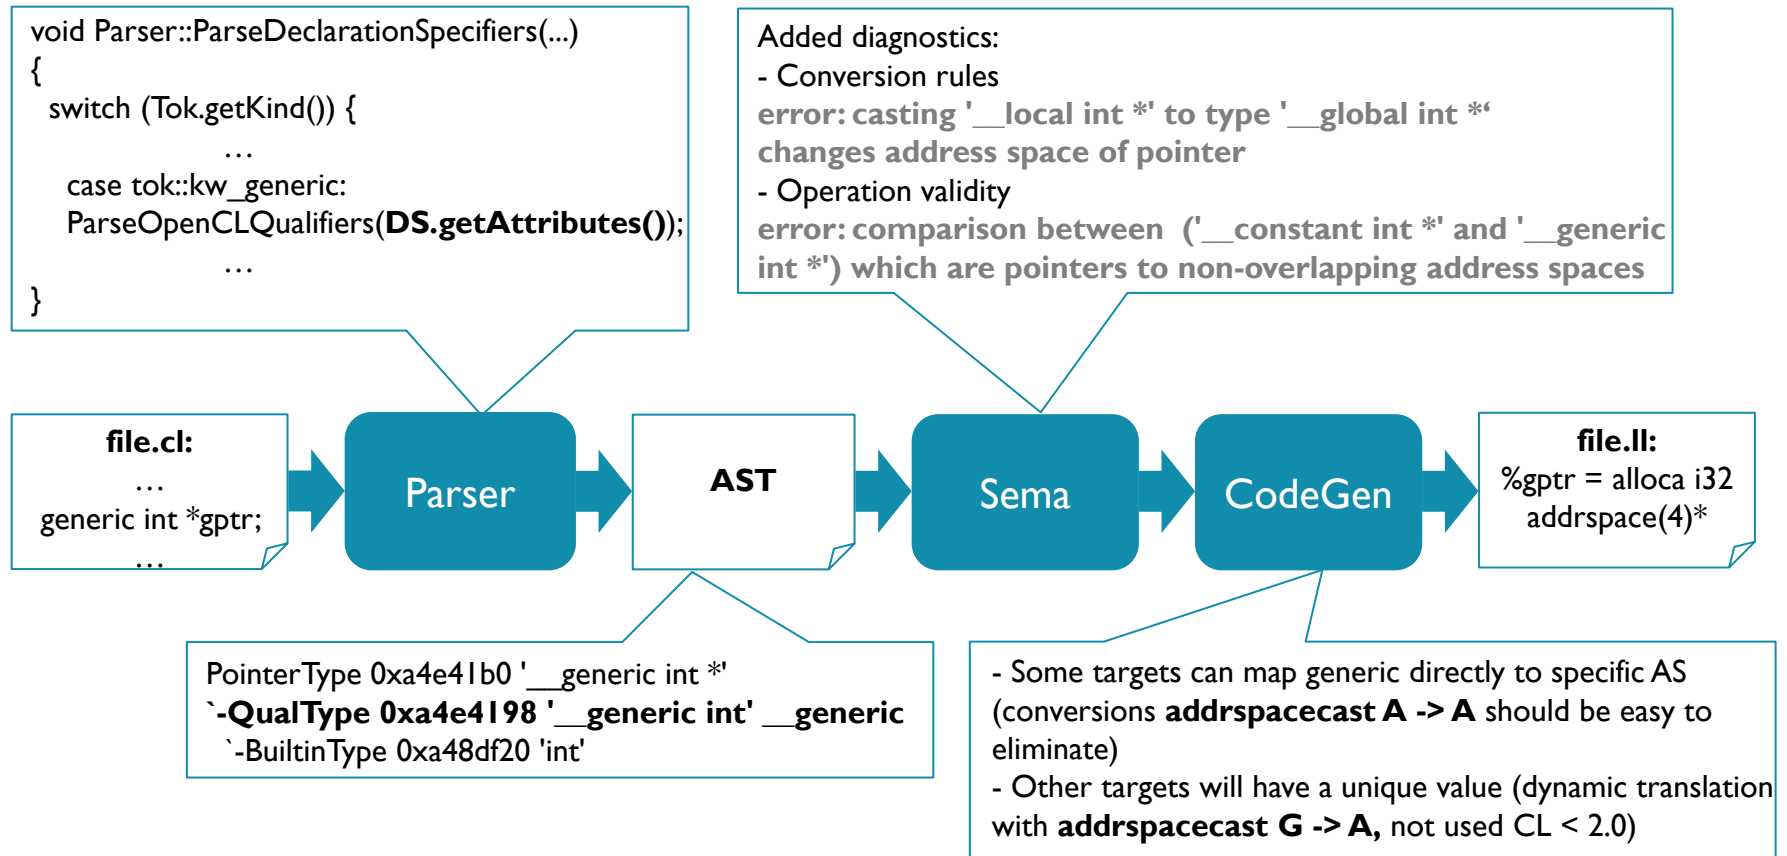
```
void foo(local int *lptr) {...}  
void foo(global int *gptr) {...}  
  
kernel void bar(local int *lptr, global int *gptr){  
    foo(lptr);  
    foo(gptr);  
}
```

OpenCL2.0

```
void foo(int *gen) {...} // only one foo is needed,  
use late binding  
  
kernel void bar(local int *lptr, global int *gptr){  
    foo(lptr); // local to generic AS conversion  
    foo(gptr); // global to generic AS conversion  
}
```


- Address Space (AS) in OpenCL is almost a part of a type
- Nothing is allowed with objects of distinct ASes including casting, operations etc.
- One of the largest changes affected Parser, Sema and CodeGen of many C paths
- Generic helps writing code more conveniently
- Easy to support in Clang reusing existing AS functionality

Generic address space in Clang



Default address space

OpenCL <2.0:



OpenCL >=2.0:

Scope Type	global	local
pointer	LangAS::generic	LangAS::generic
scalar	LangAS::global	LangAS::private

- Workable solution in order not to modify previous scheme:
 - AS is handled as a type attribute while parsing a type
 - If absent look at scope and type being parsed
 - But too early to be able to consider object kind: **NULL** - (void*)0 no AS
- We could introduce private AS explicitly as unique qualifier
 - Affects how AS is represented by previous standards
- Type printing issue (difference with the original type)

int x = &f; // warning: incompatible pointer to integer conversion initializing '**__global int**' with an expression of type ...

Atomic types

- Map CL to C11 atomic types in Clang:

Sema.cpp - Sema::Initialize():

// typedef _Atomic int atomic_int

addImplicitTypedef("atomic_int", Context.getAtomicType(Context.IntTy));

- Only subset of types are allowed
- Added Sema checks to restrict operations (only allowed through builtin functions):

atomic_int a, b; a+b; // disallowed in CL

_Atomic int a, b; a+b; // allowed in C11

- Use C11 builtin functions in Clang to implement CL2.0 functions
 - Missing memory visibility scope as LLVM doesn't have this construct

*C atomic_exchange_explicit(volatile A *obj, C desired, memory_order order, memory_scope scope); // CL*


*C atomic_exchange_explicit(volatile A *obj, C desired, memory_order order); // C11*

- Can be added as metadata or IR extension

Program scope variable

- Syntax like a global variable in C, but its value persists among different kernel executions
- Disallowed in earlier standards => Sema modification to allow
- In earlier standards we added implicit *local WG storage class* for local AS variables:
 - ***local int x;*** => Clang added *local WG storage class*
 - ***static local x;*** => Results in 2 storage classes but C allows only one
 - Removed *local WG storage* as this can be checked by an AS qualifier

Pipe



Device

```
kernel void producer(write_only pipe int p) {  
    int i = ...;  
    write_pipe(p, &i);  
}
```

```
kernel void consumer(read_only pipe int p) {  
    int i;  
    read_pipe(p, &i);  
}
```

Host

```
pipe = clCreatePipe(context, 0, sizeof(int), 10 /* # packets */...);
```

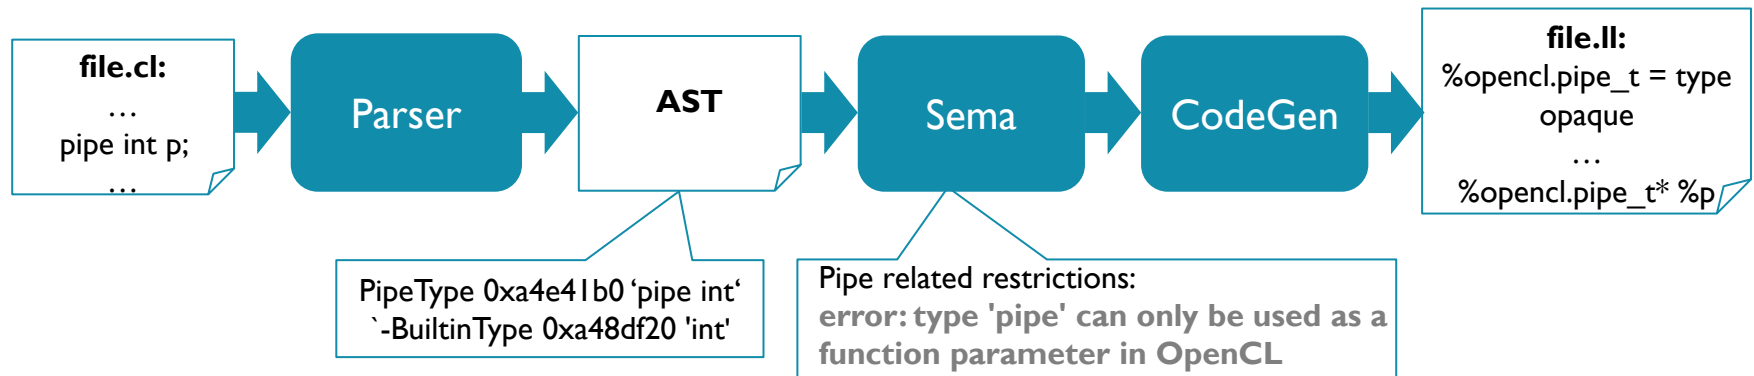
```
producer = clCreateKernel(program, "producer", &err);  
consumer = clCreateKernel(program, "consumer", &err);
```

```
err = clSetKernelArg(producer, 0, sizeof(pipe), pipe);  
err = clSetKernelArg(consumer, 0, sizeof(pipe), pipe);
```

```
err = clEnqueueNDRangeKernel(queue, producer, ...);  
err = clEnqueueNDRangeKernel(queue, consumer, ...);
```

- Classical streaming pattern
- OpenCL code specifies how elements are written/read
- Host (C/C++) code sets up pipe and connections to kernels

Pipe type



- Code repetition in Clang wrapper style types (i.e. `AtomicTypes`, `PointerTypes`, etc) and factory creation code in `ASTContext`
 - refactoring needed!
- Pipe builtin functions:
`CL: int read_pipe (read_only pipe gentype p, gentype *ptr)`
- **gentype** is any builtin or user defined type
- Generic programming style behaviour in C99
- Implemented as Clang builtin function with custom check
`Buildins.def: LANGBUILTIN(read_pipe, "i.", "tn", OCLC_LANG)`
- CodeGen to `call i32 @__read_pipe(%opencl.pipe_t* %p, i8* %ptr)`

Images

- All images are special Clang builtin types
- Handled in a similar way => a lot of copy/paste code
- OpenCL <2.0: 6 different types
 - *image1d_t, image1d_array_t, image1d_buffer_t, image2d_t, image2d_array_t, image3d_t*
- OpenCL >=2.0: 6 new types:
 - *image_2d_depth_t, image_2d_array_depth_t, image_2d_msaa_t, image_2d_array_msaa_t, image_2d_msaa_depth_t, image_2d_array_msaa_depth_t*
- Access qualifier:
 - OpenCL <2.0: *read_only/write_only*
 - OpenCL >=2.0 adds *read_write*
- Access qualifier + image type = unique type

Image problem

```
void write(write_only image2d_t img);  
  
kernel void foo(read_only image2d_t img)  
{  
    write(img); // accepted code  
}
```



```
call void @write(%opencl.image2d_t* %img);
```

But write on *write_only* is OK
In OpenCL <2.0

OpenCL 2.0 requires to call different writes
for each *write_only* and *read_write* image



- Not implemented correctly
- Access qualifiers are ignored after parsing:
 - No diagnostics wrt image access
 - No access qualifiers in IR
- Several attempts to correct
- Current review setup to correct functionality:

<http://reviews.llvm.org/D17821>

Device side enqueue

- OpenCL builtin function

```
enqueue_kernel(..., void (^block)(local void *, ...))
```

- *block* has an ObjC syntax
 - *block* can have any number of *local void** arguments
- Kind of variadic prototype
 - No standard compilation approach
 - To diagnose correctly needs to be added as Clang builtin function with a custom check

Misc features

- Loop unroll hint attribute added
 - Diagnostics and CodeGen code shared with pragma C loop hint implementation
- NOSVM attribute (but ignored)
- Still fixing AS issues in CodeGen and Sema
- Added ObjC blocks restrictions in OpenCL

`int ^bl(int, ...) = ^int(int l, ...) { // error: invalid block prototype, variadic arguments are not allowed in OpenCL`

OpenCL 2.0 current state & future work

✓ Generic AS	! Default AS	! Atomics
✓ Program scope variables	! Images	✗ Device side enqueue

- Finalise remaining work: default AS, atomics, images
- Add support for missing device side enqueue and other misc
- Improve tests and diagnostics for previous standards
- Refactoring of problematic parts

Summary

- Good progress on OpenCL2.0 (completion planned in rel3.9)
- Beneficial to derive from production quality C frontend
 - Some parts are difficult as there is no standard mechanism in Clang
 - Best use of existing C/OpenCL functionality but not affecting old functionality much
- Clang AST and internals are tailored quite well to OpenCL but IR is still very ad-hoc
 - Would it make sense to add more constructs to LLVM IR or improve support for alternative formats such as SPIR-V?

Contributors:

- ARM: Anastasia Stulova
- Intel: Alexey Bader, Xiuli PAN
- AMD: Liu Yaxun
- Tampere University of Technology: Pekka Jääskeläinen
- Others: Pedro Ferreira

Thanks! Q&A && Offline discussions

- Problematic points:
 - Unsupported targets
 - Default AS
 - Images
 - Generic IR
 - IR extensions or SPIR-V