

LLVM Libc: Status, Challenges and Future Plans

Siva Chandra Reddy
Google LLC

Agenda

1. What is LLVM libc
2. Current Status
 - a. Functions
 - b. Loader
 - c. Infrastructure
3. Challenges
4. Future Plans

LLVM libc

- A greenfield libc developed with certain goals
 - ◆ Sanitizer friendly
 - ◆ Implemented in C/C++ source code without assembly
 - ◆ Moduler

Full list of features and goals is available here:

<http://llvm.org/docs/Proposals/LLVMLibC.html>

Status: Code Metrics

- Development was kicked off a year ago
 - Over 10,000 lines of code from 200+ commits by 17 different contributors
 - Over a 100 libc functions implemented
 - Made a start to build a fully functional static-pie ELF loader
 - Build, unittest, CI and other infrastructure setup

Status: Functions

Functions from math.h

- **Basic floating point operations**
 - `fabs[f|l]`, `fdim[f|l]`, `fmax[f|l]`,
`fmin[f|l]`
- **Nearest integer functions**
 - `ceil[f|l]`, `floor[f|l]`, `trunc[f|l]`,
`round[f|l]`
- **Floating point manipulation functions**
 - `copysign[f|l]`, `frexp[f|l]`, `logb[f|l]`,
`modf[f|l]`
- **Exponential functions**
 - Single precision floating point versions of the
exponential functions `expf` and `exp2f`
- **Trigonometric functions**
 - Single precision floating point versions of
trigonometric functions `sinf`, `cosf` and `sincosf`
- **Quotient and Remainder functions**
 - `remquo[f|l]`, `remainder[f|l]`
- **Power functions**
 - `sqrt[f|l]`, `hypotf`

Functions from string.h

→ **Memory Functions:** `bzero`, `memcpy`, `memset`

- ◆ Optimized for the statistically significant subset of inputs

Function	% of calls with size ≤ 128	% of calls with size ≤ 1024
<code>memcpy</code>	96%	99%
<code>memset</code>	91%	99.9%

→ **Null terminated string functions:** `memchr`, `memrchr`,
`strcat`, `strchr`, `strcmp`, `strcpy`, `strcspn`,
`strlen`, `strnlen`, `strpbrk`, `strrchr`, `strspn`,
`strstr`, `strtok`, `strtok_r`

Status: Functions

- **Functions from threads.h**

- `call_once`, `mtx_init`, `mtx_lock`, `thrd_create`, `thrd_join` (all for Linux)
- Port them to non-linux platforms very soon.

- **Functions from signal.h**

`raise`, `sigaction`, `sigaddset`, `sigdelset`,
`sigemptyset`, `sigfillset`, `signal`, `sigprocmask` (all
for Linux)

- **Functions from ctype.h**

All `ctype.h` functions for the default locale

Miscellaneous Functions

Linux implementations of

- `abort`
- `assert`
- `errno`
- `_Exit`
- `fwrite`
- **POSIX functions** `mmap` and `munmap`

Status: Loader

Loader

- A start has been made to build a static-pie ELF loader.
- Just enough has been built to be able to test thread functions.
- Near term goal is to build a full static-pie ELF loader.

Status: Infrastructure

Build Infrastructure

- Libc specific CMake rules have been implemented
- The rules are one of the core components which make the libc implementation modular.
- Libc targets added via libc specific CMake rules have Python like fully qualified names.
 - `libc.src.math.sinf`

Status: Infrastructure

Collection of Utils

- Unit test framework
- Standalone library of C++ utilities
- Template library of floating point operations
- MPFR as reference for math function testing
 - MPFR is a C library for multi-precision floating point operations: <http://mpfr.org>

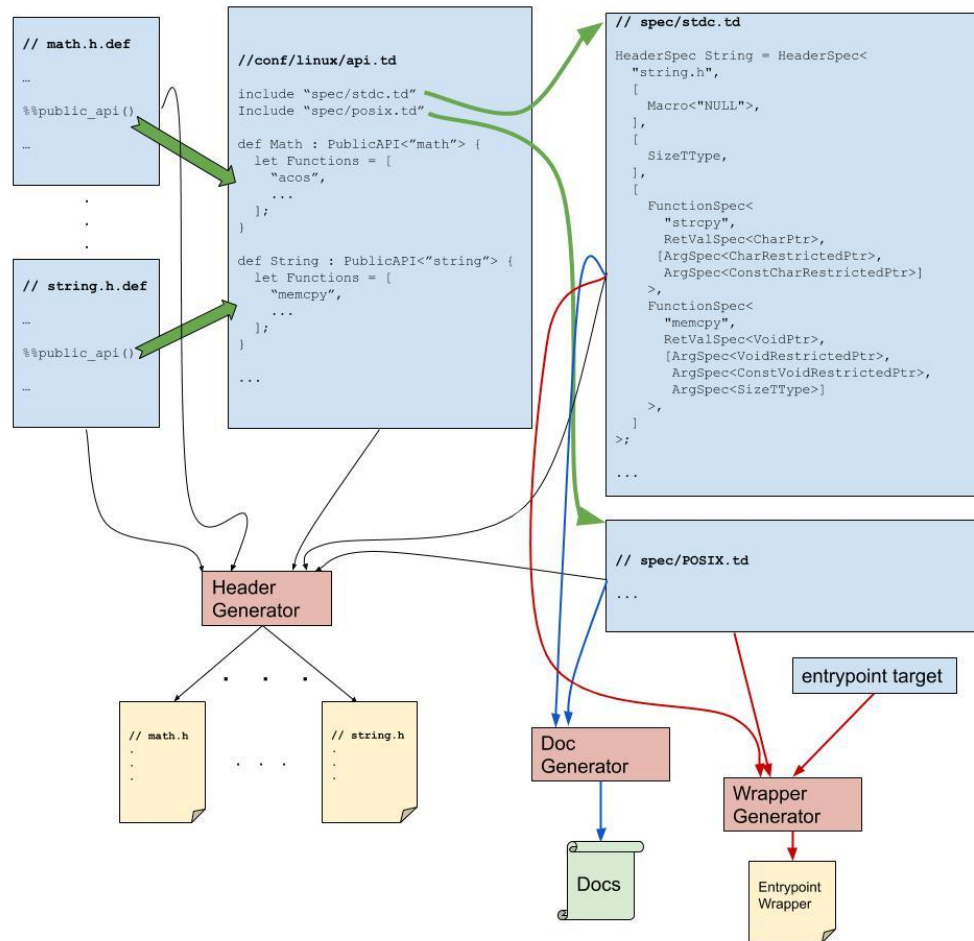
Header generation framework

- Standards are encapsulated in table gen files
- Each platform defines an API file
- The header generation tool reads the API file and generates the header files containing only the API listed in the API file
- Provides the pick the choose ability for header files

Note: A platform is a combination of the target OS + target machine architecture.

Status: Infrastructure

Header Generation Scheme



Status: Infrastructure


Clang-tidy checks

- Clang-tidy checks specific to LLVM libc have been implemented
- Protect against including undesired system headers
- Protect against polluting global namespace
- They run as part of the build and hence alert developers about deviant code at development time.
- They run on the public CI builders.

Status: Infrastructure

Redirectors

- A concept built for enabling the *mix-with-other-libcs* feature
- Redirectors are essentially wrapper functions in LLVM libc which intentionally call into the system libc
- Useful when an implementation of a particular function is not available yet in LLVM libc.



Challenges

→ **Ability to mix with other libcs**

- ◆ Challenge: Avoid header file mixup
- ◆ Challenge: Avoid symbol mixup
- ◆ Challenge: Namespace pollution

Solution: libc specific clang-tidy checks

→ **Ability to pick only parts of LLVM libc**

- ◆ Challenge: Avoid pulling in parts not required

Solution: Header generation + build system + redirectors

Future Plans

In the next one year

1. Complete the math library; that is, have an implementation for all functions coming from `math.h`.
2. Likewise, complete the null-terminated strings library.
3. Implement the API from `stdio.h` and `stdlib.h` at least for Linux.
4. Setup public CI builders for non-x86/non-linux platforms.
5. Implement full static-pie ELF loaders on linux.
6. Finish the standard threads library (`threads.h`).
7. Setup infrastructure to compare the results and performance of the math functions with similar functions from other popular libcs.
8. Run LLVM libc fuzz tests on OSS fuzz:

<https://github.com/google/oss-fuzz>

Thank You
